

O

$\text{mom}_{set} : Nx_tmpatl_{trimspaces} : n1_{set} : Nx_tmpbtl_{range} : Nnn_tmpatl_{28ifc} :$   
 $VnT_tmpbtl_{chapter_{set}} : Nx_tmpatl_{item} : Nn_tmpatl_{-1ifc} : VnTF_tmpatl_{*1433141[3]4}$

# DriverdistractionML

January 11, 2022

## 1 Detection of distracted driver using Convolutional Neural Networks

### 1.0.1 Introduction

In this project, the detection of distracted driver with Machine Learning using Convolutional Neural Networks was analysed and predicted.

### 1.0.2 Dataset description

The dataset had set of training and test images. The training data splitted into ten classes from c0 to c9. The 10 classes to predict are:

- c0: normal driving
- c1: texting - right
- c2: talking on the phone - right
- c3: texting - left
- c4: talking on the phone - left
- c5: operating the radio
- c6: drinking
- c7: reaching behind
- c8: hair and makeup
- c9: talking to passenger

Along with the set of images, two csv files were presented to assist our project. One with the details about the name of the images along with the class and another sample csv to show the submission format of the project.

### 1.0.3 Libraries

The libraries were imported to support our project. The assistance of tensorflow and keras is vital to proceed ahead. With matplotlib to plot charts and pandas to perform csv read and write operations.

```
[1]: import os
from os.path import join
import tensorflow as tf
import keras_preprocessing
from keras_preprocessing import image
from keras_preprocessing.image import ImageDataGenerator
```

```
import matplotlib.pyplot as plt
import pandas as pd
```

#### 1.0.4 Model

The Convolutional Neural Network was constructed with input size of (100,100) with the '3' represents 'rgb' format of the image. With Batch normalization, we can standardize the data in between convolutional layers. Maxpooling is to find out the maximum value from the region covered by filter and the data will be converted to one dimensional array using flatten and dropout will help us to prevent overfitting. The dense layers were added to improve efficiency and with the final dense layer represents output with 10 classes. The optimizer 'adam' was used to compile the model.

```
[2]: cnnmodel = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(32, (3,3), activation='relu', input_shape=(100, 100, 3)),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Conv2D(64, (3,3), activation='relu', padding = 'same'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Conv2D(64, (3,3), activation='relu', padding = 'same'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Conv2D(128, (3,3), activation='relu', padding = 'same'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(1024, activation='relu'),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
cnnmodel.compile(loss = 'categorical_crossentropy', optimizer = 'adam', metrics_
=> ['accuracy'])
cnnmodel.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 98, 98, 32)	896
batch_normalization (Batch Normalization)	(None, 98, 98, 32)	128
max_pooling2d (MaxPooling2D)	(None, 49, 49, 32)	0
conv2d_1 (Conv2D)	(None, 49, 49, 64)	18496
batch_normalization_1 (Batch Normalization)	(None, 49, 49, 64)	256

conv2d_2 (Conv2D)	(None, 49, 49, 64)	36928
batch_normalization_2 (Batch Normalization)	(None, 49, 49, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 24, 24, 64)	0
conv2d_3 (Conv2D)	(None, 24, 24, 128)	73856
batch_normalization_3 (Batch Normalization)	(None, 24, 24, 128)	512
flatten (Flatten)	(None, 73728)	0
dropout (Dropout)	(None, 73728)	0
dense (Dense)	(None, 1024)	75498496
dense_1 (Dense)	(None, 512)	524800
dense_2 (Dense)	(None, 10)	5130

=====  
 Total params: 76,159,754  
 Trainable params: 76,159,178  
 Non-trainable params: 576  
 =====

### 1.0.5 Data preprocessing

The working directory was set to access the folder contains training images.

```
[3]: workingdir = os.path.abspath('')
trainingdirectory = os.path.join(workingdir + '/
↳state-farm-distracted-driver-detection/imgs/train/')
```

### 1.0.6 Train and validation dataset split

Image generator was built to get access of images from the training folder. The data augmentation was used to generalize the model with horizontal flip, width and height shift range and rotation range. To improve the model, the training data were splitted into training and validation data in the ratio of 80:20. From the generator, the training and validation set can be accessed in the batch size of 64.

```
[4]: trainingdataimage = ImageDataGenerator(rescale = 1./255, rotation_range = 40,
↳width_shift_range = 0.2,
height_shift_range = 0.2, shear_range =
↳0.2, zoom_range = 0.2,
horizontal_flip = True, fill_mode =
↳'nearest', validation_split = 0.2)
trainingset = trainingdataimage.flow_from_directory(trainingdirectory,
```

```

target_size = (100,
→100), batch_size = 64,
class_mode =
→'categorical', subset = 'training', shuffle = True)
validationset = trainingdataimage.flow_from_directory(trainingdirectory,
target_size = (100,
→100), batch_size = 64,
class_mode =
→'categorical', subset = 'validation', shuffle = True)

```

Found 17943 images belonging to 10 classes.

Found 4481 images belonging to 10 classes.

### 1.0.7 Model fit

The training and validation images made to fit with the CNN model on 60 iterations with steps per iteration will be length of the image generator of training set and length of image generator of validation set was denoted as validation steps.

```

[5]: history = cnnmodel.fit(trainingset, epochs = 60, steps_per_epoch =
→len(trainingset),
validation_data = validationset, verbose = 1,
→validation_steps = len(validationset))

```

Epoch 1/60

281/281 [=====] - 114s 391ms/step - loss: 2.3757 -  
accuracy: 0.1876 - val\_loss: 9.6513 - val\_accuracy: 0.1033

Epoch 2/60

281/281 [=====] - 110s 389ms/step - loss: 1.7734 -  
accuracy: 0.3377 - val\_loss: 2.3622 - val\_accuracy: 0.2540

Epoch 3/60

281/281 [=====] - 110s 391ms/step - loss: 1.4419 -  
accuracy: 0.4602 - val\_loss: 1.5239 - val\_accuracy: 0.4300

Epoch 4/60

281/281 [=====] - 109s 389ms/step - loss: 1.1783 -  
accuracy: 0.5696 - val\_loss: 1.1394 - val\_accuracy: 0.5876

Epoch 5/60

281/281 [=====] - 110s 390ms/step - loss: 0.9663 -  
accuracy: 0.6582 - val\_loss: 1.4248 - val\_accuracy: 0.5242

Epoch 6/60

281/281 [=====] - 110s 392ms/step - loss: 0.7862 -  
accuracy: 0.7288 - val\_loss: 0.9391 - val\_accuracy: 0.6800

Epoch 7/60

281/281 [=====] - 110s 390ms/step - loss: 0.6675 -  
accuracy: 0.7702 - val\_loss: 1.0366 - val\_accuracy: 0.6637

Epoch 8/60

281/281 [=====] - 110s 392ms/step - loss: 0.5899 -  
accuracy: 0.8018 - val\_loss: 0.7342 - val\_accuracy: 0.7581

Epoch 9/60  
 281/281 [=====] - 110s 391ms/step - loss: 0.5231 - accuracy: 0.8212 - val\_loss: 0.9168 - val\_accuracy: 0.7304

Epoch 10/60  
 281/281 [=====] - 110s 390ms/step - loss: 0.4707 - accuracy: 0.8436 - val\_loss: 0.9220 - val\_accuracy: 0.7364

Epoch 11/60  
 281/281 [=====] - 110s 391ms/step - loss: 0.4250 - accuracy: 0.8570 - val\_loss: 0.6162 - val\_accuracy: 0.8045

Epoch 12/60  
 281/281 [=====] - 110s 391ms/step - loss: 0.4047 - accuracy: 0.8667 - val\_loss: 0.8444 - val\_accuracy: 0.7193

Epoch 13/60  
 281/281 [=====] - 110s 391ms/step - loss: 0.3838 - accuracy: 0.8721 - val\_loss: 0.7112 - val\_accuracy: 0.7706

Epoch 14/60  
 281/281 [=====] - 110s 393ms/step - loss: 0.3554 - accuracy: 0.8835 - val\_loss: 0.3851 - val\_accuracy: 0.8764

Epoch 15/60  
 281/281 [=====] - 111s 393ms/step - loss: 0.3435 - accuracy: 0.8886 - val\_loss: 0.3924 - val\_accuracy: 0.8744

Epoch 16/60  
 281/281 [=====] - 111s 392ms/step - loss: 0.3276 - accuracy: 0.8963 - val\_loss: 0.5650 - val\_accuracy: 0.8201

Epoch 17/60  
 281/281 [=====] - 110s 390ms/step - loss: 0.3142 - accuracy: 0.8964 - val\_loss: 0.4193 - val\_accuracy: 0.8592

Epoch 18/60  
 281/281 [=====] - 126s 449ms/step - loss: 0.3005 - accuracy: 0.9049 - val\_loss: 0.9475 - val\_accuracy: 0.7242

Epoch 19/60  
 281/281 [=====] - 113s 401ms/step - loss: 0.2899 - accuracy: 0.9072 - val\_loss: 0.4327 - val\_accuracy: 0.8596

Epoch 20/60  
 281/281 [=====] - 111s 395ms/step - loss: 0.2800 - accuracy: 0.9116 - val\_loss: 0.4845 - val\_accuracy: 0.8453

Epoch 21/60  
 281/281 [=====] - 112s 397ms/step - loss: 0.2797 - accuracy: 0.9104 - val\_loss: 0.3755 - val\_accuracy: 0.8842

Epoch 22/60  
 281/281 [=====] - 110s 392ms/step - loss: 0.2530 - accuracy: 0.9185 - val\_loss: 1.1566 - val\_accuracy: 0.6987

Epoch 23/60  
 281/281 [=====] - 111s 394ms/step - loss: 0.2626 - accuracy: 0.9151 - val\_loss: 0.6235 - val\_accuracy: 0.8286

Epoch 24/60  
 281/281 [=====] - 110s 392ms/step - loss: 0.2279 - accuracy: 0.9265 - val\_loss: 0.4207 - val\_accuracy: 0.8723

Epoch 25/60  
 281/281 [=====] - 111s 394ms/step - loss: 0.2392 - accuracy: 0.9253 - val\_loss: 0.3143 - val\_accuracy: 0.9005

Epoch 26/60  
 281/281 [=====] - 111s 393ms/step - loss: 0.2376 - accuracy: 0.9278 - val\_loss: 0.5371 - val\_accuracy: 0.8561

Epoch 27/60  
 281/281 [=====] - 111s 394ms/step - loss: 0.2296 - accuracy: 0.9249 - val\_loss: 0.4364 - val\_accuracy: 0.8788

Epoch 28/60  
 281/281 [=====] - 112s 399ms/step - loss: 0.2225 - accuracy: 0.9289 - val\_loss: 0.7858 - val\_accuracy: 0.7884

Epoch 29/60  
 281/281 [=====] - 112s 397ms/step - loss: 0.2192 - accuracy: 0.9294 - val\_loss: 0.3663 - val\_accuracy: 0.8909

Epoch 30/60  
 281/281 [=====] - 111s 395ms/step - loss: 0.2199 - accuracy: 0.9308 - val\_loss: 0.3474 - val\_accuracy: 0.8971

Epoch 31/60  
 281/281 [=====] - 111s 393ms/step - loss: 0.1921 - accuracy: 0.9390 - val\_loss: 0.5469 - val\_accuracy: 0.8422

Epoch 32/60  
 281/281 [=====] - 111s 393ms/step - loss: 0.2081 - accuracy: 0.9347 - val\_loss: 0.5347 - val\_accuracy: 0.8514

Epoch 33/60  
 281/281 [=====] - 111s 394ms/step - loss: 0.1980 - accuracy: 0.9399 - val\_loss: 0.3975 - val\_accuracy: 0.8802

Epoch 34/60  
 281/281 [=====] - 111s 393ms/step - loss: 0.1897 - accuracy: 0.9445 - val\_loss: 0.3922 - val\_accuracy: 0.8875

Epoch 35/60  
 281/281 [=====] - 111s 393ms/step - loss: 0.1963 - accuracy: 0.9390 - val\_loss: 0.3592 - val\_accuracy: 0.8947

Epoch 36/60  
 281/281 [=====] - 111s 394ms/step - loss: 0.1797 - accuracy: 0.9461 - val\_loss: 0.4619 - val\_accuracy: 0.8657

Epoch 37/60  
 281/281 [=====] - 110s 390ms/step - loss: 0.1861 - accuracy: 0.9432 - val\_loss: 0.2377 - val\_accuracy: 0.9264

Epoch 38/60  
 281/281 [=====] - 111s 395ms/step - loss: 0.1728 - accuracy: 0.9461 - val\_loss: 0.2581 - val\_accuracy: 0.9252

Epoch 39/60  
 281/281 [=====] - 111s 396ms/step - loss: 0.1836 - accuracy: 0.9458 - val\_loss: 0.2636 - val\_accuracy: 0.9154

Epoch 40/60  
 281/281 [=====] - 110s 392ms/step - loss: 0.1711 - accuracy: 0.9468 - val\_loss: 0.6902 - val\_accuracy: 0.8137

Epoch 41/60  
 281/281 [=====] - 111s 394ms/step - loss: 0.1743 - accuracy: 0.9476 - val\_loss: 0.4437 - val\_accuracy: 0.8661

Epoch 42/60  
 281/281 [=====] - 111s 393ms/step - loss: 0.1788 - accuracy: 0.9447 - val\_loss: 0.5826 - val\_accuracy: 0.8333

Epoch 43/60  
 281/281 [=====] - 110s 391ms/step - loss: 0.1790 - accuracy: 0.9473 - val\_loss: 0.3731 - val\_accuracy: 0.8822

Epoch 44/60  
 281/281 [=====] - 111s 395ms/step - loss: 0.1613 - accuracy: 0.9521 - val\_loss: 0.2974 - val\_accuracy: 0.9210

Epoch 45/60  
 281/281 [=====] - 113s 403ms/step - loss: 0.1425 - accuracy: 0.9546 - val\_loss: 0.3171 - val\_accuracy: 0.9052

Epoch 46/60  
 281/281 [=====] - 118s 418ms/step - loss: 0.1465 - accuracy: 0.9531 - val\_loss: 0.3139 - val\_accuracy: 0.9065

Epoch 47/60  
 281/281 [=====] - 116s 413ms/step - loss: 0.1480 - accuracy: 0.9529 - val\_loss: 0.2613 - val\_accuracy: 0.9165

Epoch 48/60  
 281/281 [=====] - 114s 406ms/step - loss: 0.1512 - accuracy: 0.9539 - val\_loss: 0.2615 - val\_accuracy: 0.9324

Epoch 49/60  
 281/281 [=====] - 114s 406ms/step - loss: 0.1450 - accuracy: 0.9558 - val\_loss: 0.2617 - val\_accuracy: 0.9243

Epoch 50/60  
 281/281 [=====] - 117s 415ms/step - loss: 0.1535 - accuracy: 0.9526 - val\_loss: 0.2505 - val\_accuracy: 0.9284

Epoch 51/60  
 281/281 [=====] - 116s 411ms/step - loss: 0.1430 - accuracy: 0.9564 - val\_loss: 0.2140 - val\_accuracy: 0.9344

Epoch 52/60  
 281/281 [=====] - 114s 405ms/step - loss: 0.1474 - accuracy: 0.9563 - val\_loss: 0.2797 - val\_accuracy: 0.9199

Epoch 53/60  
 281/281 [=====] - 115s 410ms/step - loss: 0.1425 - accuracy: 0.9572 - val\_loss: 0.2915 - val\_accuracy: 0.9074

Epoch 54/60  
 281/281 [=====] - 115s 409ms/step - loss: 0.1470 - accuracy: 0.9568 - val\_loss: 0.1852 - val\_accuracy: 0.9453

Epoch 55/60  
 281/281 [=====] - 115s 409ms/step - loss: 0.1298 - accuracy: 0.9589 - val\_loss: 0.1680 - val\_accuracy: 0.9543

Epoch 56/60  
 281/281 [=====] - 115s 407ms/step - loss: 0.1350 - accuracy: 0.9583 - val\_loss: 0.2776 - val\_accuracy: 0.9110



```

Epoch 57/60
281/281 [=====] - 115s 409ms/step - loss: 0.1307 -
accuracy: 0.9595 - val_loss: 0.4642 - val_accuracy: 0.9016
Epoch 58/60
281/281 [=====] - 115s 410ms/step - loss: 0.1422 -
accuracy: 0.9569 - val_loss: 0.2760 - val_accuracy: 0.9177
Epoch 59/60
281/281 [=====] - 115s 408ms/step - loss: 0.1331 -
accuracy: 0.9596 - val_loss: 0.3813 - val_accuracy: 0.9027
Epoch 60/60
281/281 [=====] - 115s 407ms/step - loss: 0.1243 -
accuracy: 0.9631 - val_loss: 0.2915 - val_accuracy: 0.9223

```

### 1.0.8 Plot to show training accuracy vs validation accuracy

The two plots were designed to visualize the learning curve of the model. One plot concentrated on Training and validation accuracy over 60 iterations and next one concentrated on training and validation loss.

```


[6]: acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']
    loss = history.history['loss']
    val_loss = history.history['val_loss']

    epochs = range(len(acc))


    plt.plot(epochs, acc, 'r', label='Training accuracy')
    plt.plot(epochs, val_acc, 'b', label='Validation accuracy')
    plt.title('Training and validation accuracy')
    plt.legend(loc=0)
    plt.figure()
    plt.show()

    plt.plot(epochs, loss, 'r', label='Training loss')
    plt.plot(epochs, val_loss, 'b', label='Validation loss')
    plt.title('Training and validation loss')
    plt.legend(loc=0)
    plt.figure()
    plt.show()

```



<Figure size 432x288 with 0 Axes>



<Figure size 432x288 with 0 Axes>

### 1.0.9 Test data prediction

Like training data, image generator were built for test data along with its directory. The important step of the project is the prediction of test images with the learning the CNN model has undergone with training and validation images.

```
[49]: testparentdirectory = os.path.join(workingdir + '/'
      ↪state-farm-distracted-driver-detection/imgs/')
testdataimage = ImageDataGenerator(rescale = 1./255)
testdata = testdataimage.flow_from_directory(testparentdirectory,
      ↪classes=['test'], target_size = (100,100))
testoutput = cnnmodel.predict(testdata, verbose = 1)
```

Found 79726 images belonging to 1 classes.

2492/2492 [=====] - 247s 99ms/step

### 1.0.10 Preparing output dataframe

The sample submission csv was read by pandas to prepare the format of output. With image and images name taken from the csv, the prediction values was replaced with the original value of csv in the same format and convert it to a dataframe to export it easily into a csv file.

```
[50]: specimencsv = pd.read_csv(os.path.join(workingdir + '/'
      ↪state-farm-distracted-driver-detection/sample_submission.csv'))
result = {'img':list(specimencsv.values[:,0]),}
for value in range(0,10):
    result['c' + str(value)] = list(testoutput[:,value])
```


```
[53]: testoutput = pd.DataFrame(result)
```

### 1.0.11 Exploratory data analysis

The number of images in each class were depicted with a bar plot.

```
[57]: imagescsv = pd.read_csv(os.path.join(workingdir + '/'
      ↪state-farm-distracted-driver-detection/driver_imgs_list.csv'))
imagescsv.classname.value_counts().plot(kind = 'bar', label = 'index')
plt.title('Number of images per class in training data')
```

```
[57]: Text(0.5, 1.0, 'Number of images per class in training data')
```



### 1.0.12 Preparing output file

Then, the csv had been written from the 'testoutput' dataframe.

```
[60]: testoutput.to_csv('Testoutput.csv', index = False, encoding='utf-8')
```

### 1.0.13 Conclusion

Thus, the prediction of test images from the model with the learning of training and validation images was successfully exported as a csv file with over 96% accuracy.