

O

$\text{mom}_{set} : Nx_tmpatl_{trimspaces} : n1_{set} : Nx_tmpbtl_{range} : Nnn_tmpatl_{28ifc} :$
 $VnT_tmpbtl_{chapter_{set}} : Nx_tmpatl_{item} : Nn_tmpatl_{-1ifc} : VnTF_tmpatl_{*1433141[3]4}$

DriverdistractionML

January 11, 2022

1 Driver distraction using Convolutional Neural Networks

1.0.1 Libraries

```
[1]: import os
from os.path import join
import tensorflow as tf
import keras_preprocessing
from keras_preprocessing import image
from keras_preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt
import pandas as pd
```

1.0.2 Model

```
[2]: cnnmodel = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(32, (3,3), activation='relu', input_shape=(100, 100, 3)),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Conv2D(64, (3,3), activation='relu', padding = 'same'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Conv2D(64, (3,3), activation='relu', padding = 'same'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Conv2D(128, (3,3), activation='relu', padding = 'same'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(1024, activation='relu'),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
cnnmodel.compile(loss = 'categorical_crossentropy', optimizer = 'adam', metrics_
    => ['accuracy'])
cnnmodel.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 98, 98, 32)	896
batch_normalization (Batch Normalization)	(None, 98, 98, 32)	128
max_pooling2d (MaxPooling2D)	(None, 49, 49, 32)	0
conv2d_1 (Conv2D)	(None, 49, 49, 64)	18496
batch_normalization_1 (Batch Normalization)	(None, 49, 49, 64)	256
conv2d_2 (Conv2D)	(None, 49, 49, 64)	36928
batch_normalization_2 (Batch Normalization)	(None, 49, 49, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 24, 24, 64)	0
conv2d_3 (Conv2D)	(None, 24, 24, 128)	73856
batch_normalization_3 (Batch Normalization)	(None, 24, 24, 128)	512
flatten (Flatten)	(None, 73728)	0
dropout (Dropout)	(None, 73728)	0
dense (Dense)	(None, 1024)	75498496
dense_1 (Dense)	(None, 512)	524800
dense_2 (Dense)	(None, 10)	5130
Total params: 76,159,754		
Trainable params: 76,159,178		
Non-trainable params: 576		

1.0.3 Data preprocessing

```
[3]: workingdir = os.path.abspath('.')
trainingdirectory = os.path.join(workingdir + '/
    ↳state-farm-distracted-driver-detection/imgs/train/')
```

1.0.4 Train and validation dataset split

```
[4]: trainingdataimage = ImageDataGenerator(rescale = 1./255, rotation_range = 40,
      width_shift_range = 0.2,
      height_shift_range = 0.2, shear_range =
      0.2, zoom_range = 0.2,
      horizontal_flip = True, fill_mode =
      'nearest', validation_split = 0.2)
trainingset = trainingdataimage.flow_from_directory(trainingdirectory,
      target_size = (100,
      100), batch_size = 64,
      class_mode =
      'categorical', subset = 'training', shuffle = True)
validationset = trainingdataimage.flow_from_directory(trainingdirectory,
      target_size = (100,
      100), batch_size = 64,
      class_mode =
      'categorical', subset = 'validation', shuffle = True)
```

Found 17943 images belonging to 10 classes.

Found 4481 images belonging to 10 classes.

1.0.5 Model fit

```
[5]: history = cnnmodel.fit(trainingset, epochs = 60, steps_per_epoch =
      len(trainingset),
      validation_data = validationset, verbose = 1,
      validation_steps = len(validationset))
```

Epoch 1/60

281/281 [=====] - 114s 391ms/step - loss: 2.3757 -
accuracy: 0.1876 - val_loss: 9.6513 - val_accuracy: 0.1033

Epoch 2/60

281/281 [=====] - 110s 389ms/step - loss: 1.7734 -
accuracy: 0.3377 - val_loss: 2.3622 - val_accuracy: 0.2540

Epoch 3/60

281/281 [=====] - 110s 391ms/step - loss: 1.4419 -
accuracy: 0.4602 - val_loss: 1.5239 - val_accuracy: 0.4300

Epoch 4/60

281/281 [=====] - 109s 389ms/step - loss: 1.1783 -
accuracy: 0.5696 - val_loss: 1.1394 - val_accuracy: 0.5876

Epoch 5/60

281/281 [=====] - 110s 390ms/step - loss: 0.9663 -
accuracy: 0.6582 - val_loss: 1.4248 - val_accuracy: 0.5242

Epoch 6/60

281/281 [=====] - 110s 392ms/step - loss: 0.7862 -
accuracy: 0.7288 - val_loss: 0.9391 - val_accuracy: 0.6800

Epoch 7/60
 281/281 [=====] - 110s 390ms/step - loss: 0.6675 - accuracy: 0.7702 - val_loss: 1.0366 - val_accuracy: 0.6637

Epoch 8/60
 281/281 [=====] - 110s 392ms/step - loss: 0.5899 - accuracy: 0.8018 - val_loss: 0.7342 - val_accuracy: 0.7581

Epoch 9/60
 281/281 [=====] - 110s 391ms/step - loss: 0.5231 - accuracy: 0.8212 - val_loss: 0.9168 - val_accuracy: 0.7304

Epoch 10/60
 281/281 [=====] - 110s 390ms/step - loss: 0.4707 - accuracy: 0.8436 - val_loss: 0.9220 - val_accuracy: 0.7364

Epoch 11/60
 281/281 [=====] - 110s 391ms/step - loss: 0.4250 - accuracy: 0.8570 - val_loss: 0.6162 - val_accuracy: 0.8045

Epoch 12/60
 281/281 [=====] - 110s 391ms/step - loss: 0.4047 - accuracy: 0.8667 - val_loss: 0.8444 - val_accuracy: 0.7193

Epoch 13/60
 281/281 [=====] - 110s 391ms/step - loss: 0.3838 - accuracy: 0.8721 - val_loss: 0.7112 - val_accuracy: 0.7706

Epoch 14/60
 281/281 [=====] - 110s 393ms/step - loss: 0.3554 - accuracy: 0.8835 - val_loss: 0.3851 - val_accuracy: 0.8764

Epoch 15/60
 281/281 [=====] - 111s 393ms/step - loss: 0.3435 - accuracy: 0.8886 - val_loss: 0.3924 - val_accuracy: 0.8744

Epoch 16/60
 281/281 [=====] - 111s 392ms/step - loss: 0.3276 - accuracy: 0.8963 - val_loss: 0.5650 - val_accuracy: 0.8201

Epoch 17/60
 281/281 [=====] - 110s 390ms/step - loss: 0.3142 - accuracy: 0.8964 - val_loss: 0.4193 - val_accuracy: 0.8592

Epoch 18/60
 281/281 [=====] - 126s 449ms/step - loss: 0.3005 - accuracy: 0.9049 - val_loss: 0.9475 - val_accuracy: 0.7242

Epoch 19/60
 281/281 [=====] - 113s 401ms/step - loss: 0.2899 - accuracy: 0.9072 - val_loss: 0.4327 - val_accuracy: 0.8596

Epoch 20/60
 281/281 [=====] - 111s 395ms/step - loss: 0.2800 - accuracy: 0.9116 - val_loss: 0.4845 - val_accuracy: 0.8453

Epoch 21/60
 281/281 [=====] - 112s 397ms/step - loss: 0.2797 - accuracy: 0.9104 - val_loss: 0.3755 - val_accuracy: 0.8842

Epoch 22/60
 281/281 [=====] - 110s 392ms/step - loss: 0.2530 - accuracy: 0.9185 - val_loss: 1.1566 - val_accuracy: 0.6987

Epoch 23/60
 281/281 [=====] - 111s 394ms/step - loss: 0.2626 - accuracy: 0.9151 - val_loss: 0.6235 - val_accuracy: 0.8286

Epoch 24/60
 281/281 [=====] - 110s 392ms/step - loss: 0.2279 - accuracy: 0.9265 - val_loss: 0.4207 - val_accuracy: 0.8723

Epoch 25/60
 281/281 [=====] - 111s 394ms/step - loss: 0.2392 - accuracy: 0.9253 - val_loss: 0.3143 - val_accuracy: 0.9005

Epoch 26/60
 281/281 [=====] - 111s 393ms/step - loss: 0.2376 - accuracy: 0.9278 - val_loss: 0.5371 - val_accuracy: 0.8561

Epoch 27/60
 281/281 [=====] - 111s 394ms/step - loss: 0.2296 - accuracy: 0.9249 - val_loss: 0.4364 - val_accuracy: 0.8788

Epoch 28/60
 281/281 [=====] - 112s 399ms/step - loss: 0.2225 - accuracy: 0.9289 - val_loss: 0.7858 - val_accuracy: 0.7884

Epoch 29/60
 281/281 [=====] - 112s 397ms/step - loss: 0.2192 - accuracy: 0.9294 - val_loss: 0.3663 - val_accuracy: 0.8909

Epoch 30/60
 281/281 [=====] - 111s 395ms/step - loss: 0.2199 - accuracy: 0.9308 - val_loss: 0.3474 - val_accuracy: 0.8971

Epoch 31/60
 281/281 [=====] - 111s 393ms/step - loss: 0.1921 - accuracy: 0.9390 - val_loss: 0.5469 - val_accuracy: 0.8422

Epoch 32/60
 281/281 [=====] - 111s 393ms/step - loss: 0.2081 - accuracy: 0.9347 - val_loss: 0.5347 - val_accuracy: 0.8514

Epoch 33/60
 281/281 [=====] - 111s 394ms/step - loss: 0.1980 - accuracy: 0.9399 - val_loss: 0.3975 - val_accuracy: 0.8802

Epoch 34/60
 281/281 [=====] - 111s 393ms/step - loss: 0.1897 - accuracy: 0.9445 - val_loss: 0.3922 - val_accuracy: 0.8875

Epoch 35/60
 281/281 [=====] - 111s 393ms/step - loss: 0.1963 - accuracy: 0.9390 - val_loss: 0.3592 - val_accuracy: 0.8947

Epoch 36/60
 281/281 [=====] - 111s 394ms/step - loss: 0.1797 - accuracy: 0.9461 - val_loss: 0.4619 - val_accuracy: 0.8657

Epoch 37/60
 281/281 [=====] - 110s 390ms/step - loss: 0.1861 - accuracy: 0.9432 - val_loss: 0.2377 - val_accuracy: 0.9264

Epoch 38/60
 281/281 [=====] - 111s 395ms/step - loss: 0.1728 - accuracy: 0.9461 - val_loss: 0.2581 - val_accuracy: 0.9252

Epoch 39/60
 281/281 [=====] - 111s 396ms/step - loss: 0.1836 - accuracy: 0.9458 - val_loss: 0.2636 - val_accuracy: 0.9154

Epoch 40/60
 281/281 [=====] - 110s 392ms/step - loss: 0.1711 - accuracy: 0.9468 - val_loss: 0.6902 - val_accuracy: 0.8137

Epoch 41/60
 281/281 [=====] - 111s 394ms/step - loss: 0.1743 - accuracy: 0.9476 - val_loss: 0.4437 - val_accuracy: 0.8661

Epoch 42/60
 281/281 [=====] - 111s 393ms/step - loss: 0.1788 - accuracy: 0.9447 - val_loss: 0.5826 - val_accuracy: 0.8333

Epoch 43/60
 281/281 [=====] - 110s 391ms/step - loss: 0.1790 - accuracy: 0.9473 - val_loss: 0.3731 - val_accuracy: 0.8822

Epoch 44/60
 281/281 [=====] - 111s 395ms/step - loss: 0.1613 - accuracy: 0.9521 - val_loss: 0.2974 - val_accuracy: 0.9210

Epoch 45/60
 281/281 [=====] - 113s 403ms/step - loss: 0.1425 - accuracy: 0.9546 - val_loss: 0.3171 - val_accuracy: 0.9052

Epoch 46/60
 281/281 [=====] - 118s 418ms/step - loss: 0.1465 - accuracy: 0.9531 - val_loss: 0.3139 - val_accuracy: 0.9065

Epoch 47/60
 281/281 [=====] - 116s 413ms/step - loss: 0.1480 - accuracy: 0.9529 - val_loss: 0.2613 - val_accuracy: 0.9165

Epoch 48/60
 281/281 [=====] - 114s 406ms/step - loss: 0.1512 - accuracy: 0.9539 - val_loss: 0.2615 - val_accuracy: 0.9324

Epoch 49/60
 281/281 [=====] - 114s 406ms/step - loss: 0.1450 - accuracy: 0.9558 - val_loss: 0.2617 - val_accuracy: 0.9243

Epoch 50/60
 281/281 [=====] - 117s 415ms/step - loss: 0.1535 - accuracy: 0.9526 - val_loss: 0.2505 - val_accuracy: 0.9284

Epoch 51/60
 281/281 [=====] - 116s 411ms/step - loss: 0.1430 - accuracy: 0.9564 - val_loss: 0.2140 - val_accuracy: 0.9344

Epoch 52/60
 281/281 [=====] - 114s 405ms/step - loss: 0.1474 - accuracy: 0.9563 - val_loss: 0.2797 - val_accuracy: 0.9199

Epoch 53/60
 281/281 [=====] - 115s 410ms/step - loss: 0.1425 - accuracy: 0.9572 - val_loss: 0.2915 - val_accuracy: 0.9074

Epoch 54/60
 281/281 [=====] - 115s 409ms/step - loss: 0.1470 - accuracy: 0.9568 - val_loss: 0.1852 - val_accuracy: 0.9453

Epoch 55/60
 281/281 [=====] - 115s 409ms/step - loss: 0.1298 - accuracy: 0.9589 - val_loss: 0.1680 - val_accuracy: 0.9543
 Epoch 56/60
 281/281 [=====] - 115s 407ms/step - loss: 0.1350 - accuracy: 0.9583 - val_loss: 0.2776 - val_accuracy: 0.9110
 Epoch 57/60
 281/281 [=====] - 115s 409ms/step - loss: 0.1307 - accuracy: 0.9595 - val_loss: 0.4642 - val_accuracy: 0.9016
 Epoch 58/60
 281/281 [=====] - 115s 410ms/step - loss: 0.1422 - accuracy: 0.9569 - val_loss: 0.2760 - val_accuracy: 0.9177
 Epoch 59/60
 281/281 [=====] - 115s 408ms/step - loss: 0.1331 - accuracy: 0.9596 - val_loss: 0.3813 - val_accuracy: 0.9027
 Epoch 60/60
 281/281 [=====] - 115s 407ms/step - loss: 0.1243 - accuracy: 0.9631 - val_loss: 0.2915 - val_accuracy: 0.9223


1.0.6 Plot to show training accuracy vs validation accuracy

```
[6]: acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']


epochs = range(len(acc))

plt.plot(epochs, acc, 'r', label='Training accuracy')
plt.plot(epochs, val_acc, 'b', label='Validation accuracy')
plt.title('Training and validation accuracy')
plt.legend(loc=0)
plt.figure()
plt.show()

plt.plot(epochs, loss, 'r', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend(loc=0)
plt.figure()
plt.show()
```

<Figure size 432x288 with 0 Axes>



<Figure size 432x288 with 0 Axes>

1.0.7 Test data prediction

```
[49]: testparentdirectory = os.path.join(workingdir + '/'
      ↪state-farm-distracted-driver-detection/imgs/')
testdataimage = ImageDataGenerator(rescale = 1./255)
testdata = testdataimage.flow_from_directory(testparentdirectory,
      ↪classes=['test'], target_size = (100,100))
testoutput = cnnmodel.predict(testdata, verbose = 1)
```

Found 79726 images belonging to 1 classes.

2492/2492 [=====] - 247s 99ms/step

1.0.8 Preparing output dataframe


```
[50]: specimencsv = pd.read_csv(os.path.join(workingdir + '/'
      ↪state-farm-distracted-driver-detection/sample_submission.csv'))
result = {'img':list(specimencsv.values[:,0]),}
for value in range(0,10):
    result['c' + str(value)] = list(testoutput[:,value])
```

```
[53]: testoutput = pd.DataFrame(result)
```

1.0.9 Exploratory data analysis

```
[57]: imagescsv = pd.read_csv(os.path.join(workingdir + '/'
      ↪state-farm-distracted-driver-detection/driver_imgs_list.csv'))
imagescsv.classname.value_counts().plot(kind = 'bar', label = 'index')
plt.title('Number of images per class in training data')
```

```
[57]: Text(0.5, 1.0, 'Number of images per class in training data')
```



1.0.10 Preparing output file

```
[60]: testoutput.to_csv('Testoutput.csv', index = False, encoding='utf-8')
```